

CPU vs GPU:

CPU

- ↳ Few cores
- ↳ Each core is very fast
- ↳ Great for sequential tasks
- ↳ Uses memory from RAM

GPU

- ↳ More cores
- ↳ Much slower
- ↳ Great for parallel tasks
- ↳ Standalone RAM

Matrix multiplication:

$$\begin{bmatrix} & & & \\ & & & \\ & & & \\ & & & \end{bmatrix}_{A \times B} \begin{bmatrix} & & & \\ & & & \\ & & & \\ & & & \end{bmatrix}_{B \times C} = \begin{bmatrix} & & & \\ & & & \\ & & & \\ & & & \end{bmatrix}_{A \times C}$$

- ↳ Compute dot products in parallel

The point of deep learning frameworks:

- 1) Easily build big computational graphs
- 2) Easily compute gradients in computational graphs
- 3) Run it all efficiently on GPU

Tensorflow:

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))
```

} input slots

```
h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
```

```
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
```

} calculations

```
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])
```

} calculate grads

```
with tf.Session() as sess:
```

```
values = {x: np.random.randn(N, D),
          w1: np.random.randn(D, H),
          w2: np.random.randn(H, D),
          y: np.random.randn(N, D),}
```

} give values

```
out = sess.run([loss, grad_w1, grad_w2],
               feed_dict=values)
loss_val, grad_w1_val, grad_w2_val = out
```

} compute
→ what to compute

Nothing is happening.
We are building
our computational
graphs

Enter tensorflow
session & do the
computation now

↳ We are copying data in between CPU & GPU / TensorFlow to numpy arrays
So,

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.Variable(tf.random_normal((D, H)))
w2 = tf.Variable(tf.random_normal((H, D)))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

learning_rate = 1e-5
new_w1 = w1.assign(w1 - learning_rate * grad_w1)
new_w2 = w2.assign(w2 - learning_rate * grad_w2)

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    for t in range(50):
        loss_val, = sess.run([loss], feed_dict=values)
```

Change w1 and w2 from **placeholder** (fed on each call) to **Variable** (persists in the graph between calls)

↳ Tells tf how we want them to be initialized

Doesn't perform cuz it doesn't need for computing Loss

Loss

→ Run once to initialize

→ Run many times to train

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.Variable(tf.random_normal((D, H)))
w2 = tf.Variable(tf.random_normal((H, D)))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

learning_rate = 1e-5
new_w1 = w1.assign(w1 - learning_rate * grad_w1)
new_w2 = w2.assign(w2 - learning_rate * grad_w2)
updates = tf.group(new_w1, new_w2)

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    losses = []
    for t in range(50):
        loss_val, _ = sess.run([loss, updates],
                               feed_dict=values)
```

hr

→ Dummy node to create dependency

Simplifies above using optimizer:

```

N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.Variable(tf.random_normal((D, H)))
w2 = tf.Variable(tf.random_normal((H, D)))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff * diff, axis=1))

optimizer = tf.train.GradientDescentOptimizer(1e-5)
updates = optimizer.minimize(loss)

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    losses = []
    for t in range(50):
        loss_val, _ = sess.run([loss, updates],
                                feed_dict=values)

```

→ Compute grads & update weights

Remember!

Compute loss using tensorflow:

`loss = tf.losses.mean_squared_error(y_pred, y)`

Higher level abstraction:

```

N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))

init = tf.contrib.layers.xavier_initializer()
h = tf.layers.dense(inputs=x, units=H,
                    activation=tf.nn.relu, kernel_initializer=init)
y_pred = tf.layers.dense(inputs=h, units=D,
                        kernel_initializer=init)

loss = tf.losses.mean_squared_error(y_pred, y)

optimizer = tf.train.GradientDescentOptimizer(1e0)
updates = optimizer.minimize(loss)

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    for t in range(50):
        loss_val, _ = sess.run([loss, updates],
                                feed_dict=values)

```

→ initializes automatically

Keras wrapper:


```
from keras.models import Sequential
from keras.layers.core import Dense, Activation
from keras.optimizers import SGD
```

```
N, D, H = 64, 1000, 100
```

```
model = Sequential()
model.add(Dense(input_dim=D, output_dim=H))
model.add(Activation('relu'))
model.add(Dense(input_dim=H, output_dim=D))
```

} make model

```
optimizer = SGD(lr=1e0)
model.compile(loss='mean_squared_error',
              optimizer=optimizer)
```

} makes graph

```
x = np.random.randn(N, D)
y = np.random.randn(N, D)
history = model.fit(x, y, nb_epoch=50,
                    batch_size=N, verbose=0)
```

} trains

Pytorch

Three layers of abstraction:

↳ Tensors:

↳ Imperative ndarray
↳ Runs on GPU

↳ Variable:

↳ Node in graph
↳ Stores variable & gradients

↳ Modules:

↳ A NN layer
↳ Store state OR learnable weights

```
import torch
dtype = torch.FloatTensor

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)

learning_rate = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()

    grad_y_pred = 2.0 * (y_pred - y)
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0
    grad_w1 = x.t().mm(grad_h)

    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2
```

For using GPU.
torch.cuda.

Autograd:

```
import torch
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10
x = Variable(torch.randn(N, D_in), requires_grad=False)
y = Variable(torch.randn(N, D_out), requires_grad=False)
w1 = Variable(torch.randn(D_in, H), requires_grad=True)
w2 = Variable(torch.randn(H, D_out), requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    if w1.grad: w1.grad.data.zero_()
    if w2.grad: w2.grad.data.zero_()
    loss.backward()

    w1.data -= learning_rate * w1.grad.data
    w2.data -= learning_rate * w2.grad.data
```

Don't grads
Need grads
Backward pass
update step

x.data is Tensor
x.grad is a Var of grads
x.grad.data is a Tensor

Custom autograd functions:

```
class ReLU(torch.autograd.Function):
    def forward(self, x):
        self.save_for_backward(x)
        return x.clamp(min=0)

    def backward(self, grad_y):
        x, = self.saved_tensors
        grad_input = grad_y.clone()
        grad_input[x < 0] = 0
        return grad_input
```

Higher level abstraction: nn

```
import torch
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10
x = Variable(torch.randn(N, D_in))
y = Variable(torch.randn(N, D_out), requires_grad=False)

model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out))
loss_fn = torch.nn.MSELoss(size_average=False)

learning_rate = 1e-4
for t in range(500):
    y_pred = model(x)
    loss = loss_fn(y_pred, y)

    model.zero_grad()
    loss.backward()

    for param in model.parameters():
        param.data -= learning_rate * param.grad.data
```

layers
loss
update step
Train

For automatic updates:

optimizer = torch.optim.Adam(model.
parameters(), lr=learning_rate)
optimizer.step() update step


Minibatches:

loader = DataLoader(TensorDataset(x,y), batch_size=8)

Static VS dynamic graph:

↳ TensorFlow:

- ↳ Build graph
- ↳ Run many times
- ↳ Optimize graphs



↳ You can serialize the graph & run it without the code that built the graph

↳ less clean

↳ PyTorch:

- ↳ Each forward pass builds a new graph
- ↳ Can't optimize graph
- ↳ Graph building & execution intertwined
- ↳ Makes code cleaner

```
N, D, H = 3, 4, 5

x = Variable(torch.randn(N, D))
w1 = Variable(torch.randn(D, H))
w2 = Variable(torch.randn(D, H))

z = 10
if z > 0:
    y = x.mm(w1)
else:
    y = x.mm(w2)
```

} easily using python code

```
N, D, H = 3, 4, 5
x = tf.placeholder(tf.float32, shape=(N, D))
z = tf.placeholder(tf.float32, shape=None)
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(D, H))
```

```
def f1(): return tf.matmul(x, w1)
def f2(): return tf.matmul(x, w2)
y = tf.cond(tf.less(z, 0), f1, f2)
```

```
with tf.Session() as sess:
```

```
    values = {
        x: np.random.randn(N, D),
        z: 10,
        w1: np.random.randn(D, H),
        w2: np.random.randn(D, H),
    }
```

```
    y_val = sess.run(y, feed_dict=values)
```

→ Need to baked into the graph